

```
import pandas as pd
import numpy as np
from tqdm import tqdm
from tqdm.notebook import tqdm_notebook
tqdm_notebook.pandas()
import warnings
warnings.filterwarnings('ignore')
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

```
! cp '/content/drive/My Drive/tweet-sentiment-extraction/preprocessed_train.csv' .
! cp '/content/drive/My Drive/tweet-sentiment-extraction/preprocessed_test.csv' .
```

```
train_df = pd.read_csv('preprocessed_train.csv')
test_df = pd.read_csv('preprocessed_test.csv')
```

```
train_df.shape, test_df.shape
```

```
↳ ((27469, 7), (3534, 3))
```

```
train_df.sample(5)
```

	textID	text	selected_text	sentiment	misspelled	start_indices	end_indices
	2736	9d4b1e13a3 jen lol i know seems that the whole fabric of ...	i know seems that the whole fabric of our eart...	neutral	No	2	22
	14160	aa04ff026b i ca not get to bed	i ca not get to bed	neutral	No	0	5
	20797	34fc124590 final thought for the day does deodorant reall...	final thought for the day does deodorant reall...	neutral	No	0	19
	3315	3c4f26ae13 well no phone today we are waiting till june t...	good	positive	No	18	18
	966	72bafdda6a i was born there	i was born there	neutral	No	0	3

```
test_df.sample(5)
```

	textID	text	sentiment
	1483	9932b9c033 if i wasnt workin in hours id be gettin ratars...	negative
	2323	166332e48a i am oh so very bored buut almost days til i l...	negative
	1061	afa094db6a got alot of runnin around to do today to get t...	positive
	1280	990c3cf35f too sick for rigging tomorrow	negative
	1261	c69ccdfa56 i am very much in tune with your words today t...	positive

```
train_df[train_df.end_indices<train_df.start_indices]
```

	textID	text	selected_text	sentiment	misspelled	start_indices	end_indices
	6393	ddbce5f751 this is great i just found out that it is star...	amay the be it	positive	No	13	8
	13668	8607d4de1a dans public transport again and have decided i...	utter curse that	negative	No	16	14

```
train_df = train_df[train_df.end_indices>=train_df.start_indices]
```

```
train_df = train_df[train_df.sentiment!='neutral']
```

```
train_df.shape
```

```
↳ (16351, 7)
```

```
X = train_df[['text', 'selected_text', 'sentiment', 'start_indices', 'end_indices']]
```

```
X.shape
```

```
↳ (16351, 5)
```

```
lens=[]
for each in X.text.values:
    lens.append(len(each.split()))
```

```
print('max length of sentence:',max(lens))
```

```
↳ max length of sentence: 33
```

For each input text, we are gonna create a output vector in such a way that, the words which are part of selected text will be given a value of 1 and others will be given a value of 0

Example : text -----> 'I am not happy with the kind of service'

selected_text--> 'not happy'

output -----> 0 0 1 1 0 0 0 0

Since the max length of input sentences are 32, output vector will be a 32 dimensional vector

```
Y = np.zeros((X.shape[0],max(lens)+1))
for i,each in tqdm(enumerate(X.values)):
    start = each[3]
    end = each[4]
    Y[i][start:end+1] = 1
```

```
↳ 16351it [00:00, 599301.48it/s]
```

```
#Cross checking whether the code has worked correctly.
import random
for _ in range(5):
    x = random.randint(0,train_df.shape[0])
    print('Data:',X.values[x])
    print('o/p vector:',Y[x])
```



```
#https://machinelearningmastery.com/use-word-embedding-layers-deep-learning-keras/
'''from numpy import asarray
from numpy import zeros
embeddings_index = dict()
with open('/content/cc.en.300.vec') as f:
#with open('/content/glove.6B.300d.txt') as f:
    for line in f:
        values = line.split()
        word = values[0]
        coefs = asarray(values[1:], dtype='float32')
        embeddings_index[word] = coefs

print('Loaded %s word vectors.' % len(embeddings_index))'''

↳ 'from numpy import asarray\nfrom numpy import zeros\nembeddings_index = dict()\nwith open('/content/cc.en.300.vec') as f:\nwith open('/content/glove.6B.300d.txt') as f:\n    for line in f:\n        values = line.split()\n        word = values[0]\n        coefs = asarray(values[1:], dtype='float32')\n        embeddings_index[word] = coefs\n        \nprint('Loaded %s word vectors.' % len(embe\n        ddings_index))'

embedding_matrix = np.zeros((vocab_size_text, 300))
for word, i in tokenizer_text.word_index.items():
    try:
        embedding_vector = fasttext_model.wv.word_vec(word)
        embedding_matrix[i] = embedding_vector
    except:
        continue
print(embedding_matrix.shape)

↳ (17397, 300)

max_length_text=33
from tensorflow.keras.preprocessing.sequence import pad_sequences
train_text = pad_sequences(train_text,maxlen=max_length_text,padding='post')
val_text = pad_sequences(val_text,maxlen=max_length_text,padding='post')
print(train_text.shape,val_text.shape)

↳ (13080, 33) (3271, 33)

train_sentiment = x_train['sentiment'].values
val_sentiment = x_val['sentiment'].values

from tensorflow.keras.preprocessing.text import Tokenizer
tokenizer_sentiment = Tokenizer(lower=True,split=' ',filters='!"#%&()*+,-./:;<=>?@[\\]^_{}~\t\n',oov_token='oov')
tokenizer_sentiment.fit_on_texts(train_sentiment)
train_sentiment=tokenizer_sentiment.texts_to_sequences(train_sentiment)
val_sentiment=tokenizer_sentiment.texts_to_sequences(val_sentiment)
print(len(train_sentiment),len(val_sentiment))
print(tokenizer_sentiment.word_index)
vocab_size_sentiment=len(tokenizer_sentiment.word_index)+1
print(vocab_size_sentiment)

↳ 13080 3271
{'oov': 1, 'positive': 2, 'negative': 3}
4

max_length_sentiment=1
from tensorflow.keras.preprocessing.sequence import pad_sequences
train_sentiment = pad_sequences(train_sentiment,maxlen=max_length_sentiment,padding='post')
val_sentiment = pad_sequences(val_sentiment,maxlen=max_length_sentiment,padding='post')
print(train_sentiment.shape,val_sentiment.shape)

↳ (13080, 1) (3271, 1)

embedding_matrix1 = np.zeros((vocab_size_sentiment, 300))
for word, i in tokenizer_sentiment.word_index.items():
    try:
        embedding_vector = fasttext_model.wv.word_vec(word)
        embedding_matrix1[i] = embedding_vector
    except:
        continue
print(embedding_matrix1.shape)

↳ (4, 300)
```

#https://machinelearningmastery.com/timedistributed-layer-for-long-short-term-memory-networks-in-python/

```
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Embedding,Dense,Dropout,Concatenate,Flatten,TimeDistributed,Input,GRU,BatchNormalization,Bidirectional,SpatialDropout1D,LSTM,Layer
from tensorflow.keras.regularizers import l2

input1=Input(shape=(max_length_text,),name='input_text')
input2=Input(shape=(max_length_sentiment,),name='input_sentiment')

embed1 = Embedding(vocab_size_text,300,input_length=max_length_text,name='embedding1',\
                    trainable=False,mask_zero = True,embeddings_initializer=tf.constant_initializer(embedding_matrix))(input1)
embed2 = Embedding(vocab_size_sentiment,300,input_length=max_length_text,name='embedding2',\
                    trainable=False,mask_zero = True,embeddings_initializer=tf.constant_initializer(embedding_matrix1))(input2)
concat= Concatenate(axis=1)([embed1,embed2])
gru=Bidirectional(GRU(8,name='gru',return_sequences=True,dropout=0.4))(concat)

import tensorflow as tf
from tensorflow.keras.activations import tanh
from tensorflow.keras.backend import dot
from tensorflow.keras.activations import softmax
from tensorflow.keras.backend import sum
class attention(tf.keras.layers.Layer):

    def __init__(self, return_sequences=True):
        self.return_sequences = return_sequences
        super().__init__()

    def build(self, input_shape):

        self.W=self.add_weight(name="att_weight", shape=(input_shape[-1],1),
                                initializer="normal")
        self.b=self.add_weight(name="att_bias", shape=(input_shape[1],1),
                                initializer="zeros")

https://colab.research.google.com/drive/1ZKgkV4JFGRc0tU-J2V-qf51uxP2z5bcf#scrollTo=wMr9w4PUwo8v&printMode=true
```

```
super(attention,self).build(input_shape)

def call(self, x):

    e = tanh(dot(x,self.W)+self.b)
    a = tf.nn.softmax(e, axis=1)
    output = x*a

    if self.return_sequences:
        return output

    return tf.reduce_sum(output, axis=1)

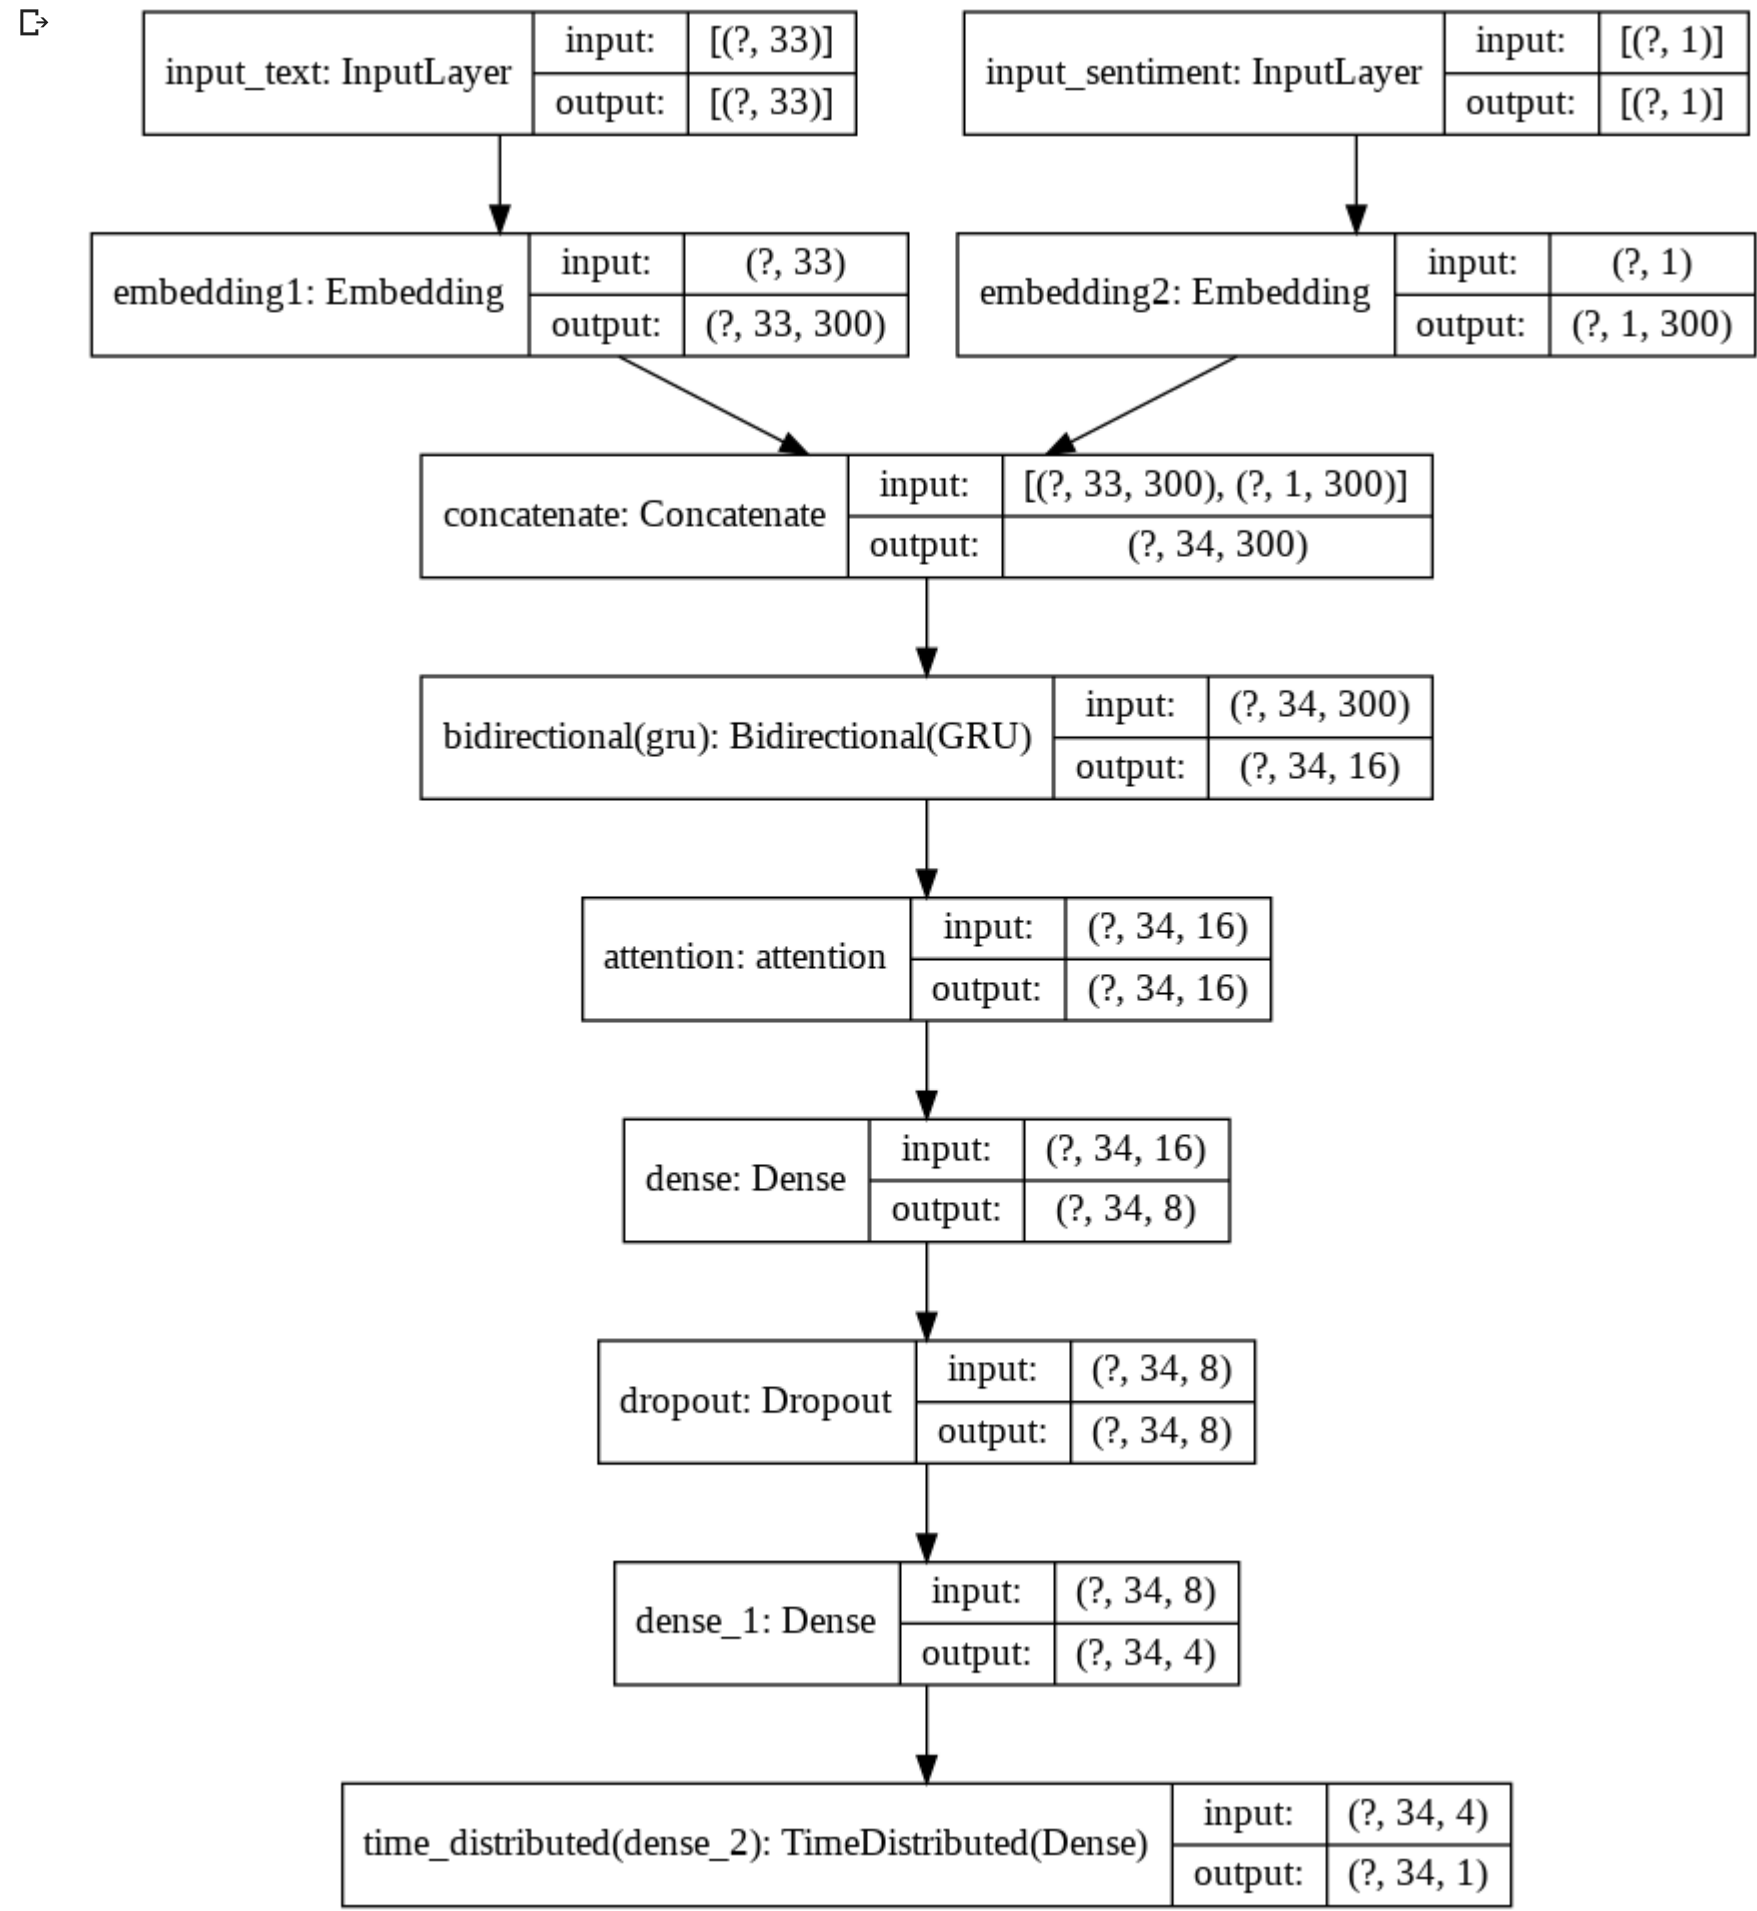
#input1=Input(shape=(max_length_text,),name='input_text')
#input2=Input(shape=(max_length_sentiment,),name='input_sentiment')
#concat= Concatenate()([input1,input2])
#embed = Embedding(vocab_size_text,300,input_length=max_length_text,name='embedding',\
#                  mask_zero = True,embeddings_initializer=tf.constant_initializer(embedding_matrix))(concat)
#gru=Bidirectional(GRU(16,name='gru',return_sequences=True,dropout=0.4))(embed)

att = attention(return_sequences=True)(gru)
dense1 = Dense(8,activation='relu',kernel_regularizer=l2(0.0001))(att)
dp = Dropout(0.5)(dense1)
dense1 = Dense(4,activation='relu',kernel_regularizer=l2(0.0001))(dp)
output=TimeDistributed(Dense(1,activation='sigmoid'))(dense1)

model=Model(inputs=[input1,input2],outputs=[output])

for each in model.layers:
    if(type(each) == tf.keras.layers.Embedding):
        each.trainable = False

import tensorflow as tf
tf.keras.utils.plot_model(model, 'Model.png', show_shapes=True)
```



```
model.summary()
```



```
Model: "functional_1"

Layer (type)                 Output Shape                 Param #   Connected to
=====
input_text (InputLayer)      [(None, 33)]                0
input_sentiment (InputLayer) [(None, 1)]                 0
embedding1 (Embedding)       (None, 33, 300)            5219100   input_text[0][0]
dropout1 (Dropout)           (None, 33, 300)            0
dense1 (Dense)               (None, 33, 8)              2640      dropout1[0][0]
dense2 (Dense)               (None, 33, 1)              34        dense1[0][0]
time_distributed1 (TimeDistributed) (None, 33, 1)              34        dense2[0][0]

%load_ext tensorboard
import datetime
import os
log_dir= os.path.join("tensorboard_logs1" , datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1, write_graph=True)
! mkdir  'checkpoint'
file_path = os.path.join('checkpoint/model2.hdf5')
ckpt_save = tf.keras.callbacks.ModelCheckpoint(filepath=file_path,save_weights_only=True,monitor='val_loss',save_best_only=True,verbose=1)
callbacks=[tensorboard_callback,ckpt_save]

def my_loss(true,pred):
    #print(true.shape,pred.shape)

    loss_obj = tf.keras.losses.BinaryCrossentropy(reduction=tf.keras.losses.Reduction.SUM)
    loss = loss_obj (true,pred)
    return loss/128 #batch size

#loss_fn = tf.keras.losses.BinaryCrossentropy()
model.compile(optimizer='adam',loss=my_loss,metrics=[ 'accuracy'])


model.fit(input_data,output_data,epochs=50,batch_size=128,validation_data=val_data,callbacks=callbacks)
```

↗

```
WARNING:tensorflow:Model failed to serialize as JSON. Ignoring... Layer attention has arguments in `__init__` and therefore must override `get_config`.
Epoch 1/50
 1/103 [.....] - ETA: 0s - loss: 23.5665 - accuracy: 0.8389WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary_ops_v2.py
Instructions for updating:
use `tf.profiler.experimental.stop` instead.
 2/103 [.....] - ETA: 18s - loss: 23.5518 - accuracy: 0.8622WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch ti
102/103 [=====>.] - ETA: 0s - loss: 21.4441 - accuracy: 0.8903
Epoch 00001: val_loss improved from inf to 18.45364, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 131ms/step - loss: 21.4115 - accuracy: 0.8903 - val_loss: 18.4536 - val_accuracy: 0.8895
Epoch 2/50
102/103 [=====>.] - ETA: 0s - loss: 16.0953 - accuracy: 0.8908
Epoch 00002: val_loss improved from 18.45364 to 13.44885, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 120ms/step - loss: 16.0703 - accuracy: 0.8908 - val_loss: 13.4488 - val_accuracy: 0.8895
Epoch 3/50
102/103 [=====>.] - ETA: 0s - loss: 12.3556 - accuracy: 0.8908
Epoch 00003: val_loss improved from 13.44885 to 10.13099, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 115ms/step - loss: 12.3362 - accuracy: 0.8908 - val_loss: 10.1310 - val_accuracy: 0.8895
Epoch 4/50
102/103 [=====>.] - ETA: 0s - loss: 9.6318 - accuracy: 0.8908
Epoch 00004: val_loss improved from 10.13099 to 8.42521, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 118ms/step - loss: 9.6173 - accuracy: 0.8908 - val_loss: 8.4252 - val_accuracy: 0.8895
Epoch 5/50
102/103 [=====>.] - ETA: 0s - loss: 8.7182 - accuracy: 0.8908
Epoch 00005: val_loss improved from 8.42521 to 8.01204, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 117ms/step - loss: 8.7049 - accuracy: 0.8908 - val_loss: 8.0120 - val_accuracy: 0.8895
Epoch 6/50
102/103 [=====>.] - ETA: 0s - loss: 8.3851 - accuracy: 0.8908
Epoch 00006: val_loss improved from 8.01204 to 7.82283, saving model to checkpoint/model12.hdf5
103/103 [=====] - 11s 103ms/step - loss: 8.3729 - accuracy: 0.8908 - val_loss: 7.8228 - val_accuracy: 0.8895
Epoch 7/50
102/103 [=====>.] - ETA: 0s - loss: 8.1768 - accuracy: 0.8908
Epoch 00007: val_loss improved from 7.82283 to 7.65312, saving model to checkpoint/model12.hdf5
103/103 [=====] - 10s 94ms/step - loss: 8.1646 - accuracy: 0.8908 - val_loss: 7.6531 - val_accuracy: 0.8895
Epoch 8/50
102/103 [=====>.] - ETA: 0s - loss: 8.0112 - accuracy: 0.8907
Epoch 00008: val_loss improved from 7.65312 to 7.52960, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 113ms/step - loss: 7.9989 - accuracy: 0.8908 - val_loss: 7.5296 - val_accuracy: 0.8895
Epoch 9/50
102/103 [=====>.] - ETA: 0s - loss: 7.9186 - accuracy: 0.8907
Epoch 00009: val_loss improved from 7.52960 to 7.46558, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 115ms/step - loss: 7.9065 - accuracy: 0.8908 - val_loss: 7.4656 - val_accuracy: 0.8895
Epoch 10/50
102/103 [=====>.] - ETA: 0s - loss: 7.8336 - accuracy: 0.8908
Epoch 00010: val_loss improved from 7.46558 to 7.41664, saving model to checkpoint/model12.hdf5
103/103 [=====] - 10s 99ms/step - loss: 7.8219 - accuracy: 0.8908 - val_loss: 7.4166 - val_accuracy: 0.8895
Epoch 11/50
102/103 [=====>.] - ETA: 0s - loss: 7.7533 - accuracy: 0.8909
Epoch 00011: val_loss improved from 7.41664 to 7.36623, saving model to checkpoint/model12.hdf5
103/103 [=====] - 10s 96ms/step - loss: 7.7423 - accuracy: 0.8908 - val_loss: 7.3662 - val_accuracy: 0.8895
Epoch 12/50
102/103 [=====>.] - ETA: 0s - loss: 7.7021 - accuracy: 0.8908
Epoch 00012: val_loss improved from 7.36623 to 7.32386, saving model to checkpoint/model12.hdf5
103/103 [=====] - 10s 94ms/step - loss: 7.6902 - accuracy: 0.8908 - val_loss: 7.3239 - val_accuracy: 0.8895
Epoch 13/50
102/103 [=====>.] - ETA: 0s - loss: 7.6458 - accuracy: 0.8907
Epoch 00013: val_loss improved from 7.32386 to 7.30067, saving model to checkpoint/model12.hdf5
103/103 [=====] - 11s 108ms/step - loss: 7.6341 - accuracy: 0.8908 - val_loss: 7.3007 - val_accuracy: 0.8895
Epoch 14/50
102/103 [=====>.] - ETA: 0s - loss: 7.5804 - accuracy: 0.8909
Epoch 00014: val_loss improved from 7.30067 to 7.27068, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 115ms/step - loss: 7.5701 - accuracy: 0.8908 - val_loss: 7.2707 - val_accuracy: 0.8895
Epoch 15/50
102/103 [=====>.] - ETA: 0s - loss: 7.5668 - accuracy: 0.8908
Epoch 00015: val_loss improved from 7.27068 to 7.25156, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 116ms/step - loss: 7.5553 - accuracy: 0.8908 - val_loss: 7.2516 - val_accuracy: 0.8895
Epoch 16/50
102/103 [=====>.] - ETA: 0s - loss: 7.5502 - accuracy: 0.8908
Epoch 00016: val_loss improved from 7.25156 to 7.23627, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 116ms/step - loss: 7.5390 - accuracy: 0.8908 - val_loss: 7.2363 - val_accuracy: 0.8895
Epoch 17/50
102/103 [=====>.] - ETA: 0s - loss: 7.5178 - accuracy: 0.8908
Epoch 00017: val_loss improved from 7.23627 to 7.22623, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 115ms/step - loss: 7.5068 - accuracy: 0.8908 - val_loss: 7.2262 - val_accuracy: 0.8895
Epoch 18/50
102/103 [=====>.] - ETA: 0s - loss: 7.5015 - accuracy: 0.8907
Epoch 00018: val_loss improved from 7.22623 to 7.21520, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 112ms/step - loss: 7.4898 - accuracy: 0.8908 - val_loss: 7.2152 - val_accuracy: 0.8895
Epoch 19/50
102/103 [=====>.] - ETA: 0s - loss: 7.4870 - accuracy: 0.8908
Epoch 00019: val_loss did not improve from 7.21520
103/103 [=====] - 11s 112ms/step - loss: 7.4758 - accuracy: 0.8908 - val_loss: 7.2293 - val_accuracy: 0.8895
Epoch 20/50
102/103 [=====>.] - ETA: 0s - loss: 7.4668 - accuracy: 0.8908
Epoch 00020: val_loss improved from 7.21520 to 7.19120, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 115ms/step - loss: 7.4553 - accuracy: 0.8908 - val_loss: 7.1912 - val_accuracy: 0.8895
Epoch 21/50
102/103 [=====>.] - ETA: 0s - loss: 7.4587 - accuracy: 0.8907
Epoch 00021: val_loss improved from 7.19120 to 7.19000, saving model to checkpoint/model12.hdf5
103/103 [=====] - 11s 110ms/step - loss: 7.4468 - accuracy: 0.8908 - val_loss: 7.1900 - val_accuracy: 0.8895
Epoch 22/50
102/103 [=====>.] - ETA: 0s - loss: 7.4400 - accuracy: 0.8907
Epoch 00022: val_loss improved from 7.19000 to 7.18933, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 112ms/step - loss: 7.4284 - accuracy: 0.8908 - val_loss: 7.1893 - val_accuracy: 0.8895
Epoch 23/50
102/103 [=====>.] - ETA: 0s - loss: 7.4234 - accuracy: 0.8908
Epoch 00023: val_loss improved from 7.18933 to 7.17182, saving model to checkpoint/model12.hdf5
103/103 [=====] - 14s 137ms/step - loss: 7.4127 - accuracy: 0.8908 - val_loss: 7.1718 - val_accuracy: 0.8895
Epoch 24/50
102/103 [=====>.] - ETA: 0s - loss: 7.4248 - accuracy: 0.8909
Epoch 00024: val_loss improved from 7.17182 to 7.16440, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 123ms/step - loss: 7.4141 - accuracy: 0.8908 - val_loss: 7.1644 - val_accuracy: 0.8895
Epoch 25/50
102/103 [=====>.] - ETA: 0s - loss: 7.4063 - accuracy: 0.8907
Epoch 00025: val_loss did not improve from 7.16440
103/103 [=====] - 11s 108ms/step - loss: 7.3947 - accuracy: 0.8908 - val_loss: 7.1827 - val_accuracy: 0.8895
Epoch 26/50
102/103 [=====>.] - ETA: 0s - loss: 7.3948 - accuracy: 0.8909
Epoch 00026: val_loss improved from 7.16440 to 7.14186, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 119ms/step - loss: 7.3841 - accuracy: 0.8908 - val_loss: 7.1419 - val_accuracy: 0.8895
Epoch 27/50
102/103 [=====>.] - ETA: 0s - loss: 7.3936 - accuracy: 0.8907
Epoch 00027: val_loss improved from 7.14186 to 7.13866, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 123ms/step - loss: 7.3821 - accuracy: 0.8908 - val_loss: 7.1387 - val_accuracy: 0.8895
Epoch 28/50
102/103 [=====>.] - ETA: 0s - loss: 7.3861 - accuracy: 0.8908
Epoch 00028: val_loss improved from 7.13866 to 7.13403, saving model to checkpoint/model12.hdf5
103/103 [=====] - 12s 119ms/step - loss: 7.3748 - accuracy: 0.8908 - val_loss: 7.1340 - val_accuracy: 0.8895
Epoch 29/50
102/103 [=====>.] - ETA: 0s - loss: 7.3689 - accuracy: 0.8907
Epoch 00029: val_loss did not improve from 7.13403
103/103 [=====] - 13s 124ms/step - loss: 7.3575 - accuracy: 0.8908 - val_loss: 7.1442 - val_accuracy: 0.8895
Epoch 30/50
102/103 [=====>.] - ETA: 0s - loss: 7.3573 - accuracy: 0.8907
Epoch 00030: val_loss did not improve from 7.13403
103/103 [=====] - 13s 124ms/step - loss: 7.3460 - accuracy: 0.8908 - val_loss: 7.1474 - val_accuracy: 0.8895
Epoch 31/50
102/103 [=====>.] - ETA: 0s - loss: 7.3558 - accuracy: 0.8941
```

Epoch 00031: val_loss improved from 7.13403 to 7.11304, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 122ms/step - loss: 7.3450 - accuracy: 0.8941 - val_loss: 7.1130 - val_accuracy: 0.9019
Epoch 32/50
102/103 [=====>.] - ETA: 0s - loss: 7.3510 - accuracy: 0.8996
Epoch 00032: val_loss improved from 7.11304 to 7.11153, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 124ms/step - loss: 7.3401 - accuracy: 0.8995 - val_loss: 7.1115 - val_accuracy: 0.9015
Epoch 33/50
102/103 [=====>.] - ETA: 0s - loss: 7.3295 - accuracy: 0.8997
Epoch 00033: val_loss improved from 7.11153 to 7.10574, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 126ms/step - loss: 7.3186 - accuracy: 0.8997 - val_loss: 7.1057 - val_accuracy: 0.9020
Epoch 34/50
102/103 [=====>.] - ETA: 0s - loss: 7.3121 - accuracy: 0.9000
Epoch 00034: val_loss improved from 7.10574 to 7.09552, saving model to checkpoint/model12.hdf5
103/103 [=====] - 13s 126ms/step - loss: 7.3014 - accuracy: 0.8999 - val_loss: 7.0955 - val_accuracy: 0.9017
Epoch 35/50
102/103 [=====>.] - ETA: 0s - loss: 7.3233 - accuracy: 0.8997
Epoch 00035: val loss did not improve from 7.09552

tf.keras.backend.clear_session()
%tensorboard --logdir \$log_dir --port 0



Epoch 00040: val_loss improved from 7.05430 to 7.05030, saving model to checkpoint/model12.hdf5
102/103 [=====>.] - ETA: 0s - loss: 7.2032 - accuracy: 0.9002

```
#input_data = (train_text,train_sentiment)
#output_data = y_train
#val = (val_text,val_sentiment)
#output_val = y_val
#val_data = (val,output_val)
```

Inference

model.load_weights('checkpoint/model12.hdf5')

w = model.get_weights()
len(w)

16

names = [weight.name for layer in model.layers for weight in layer.weights]
len(names)

16

for layer,weight in zip(names,w):
 print(layer,weight.shape)

embedding1/embeddings:0 (17397, 300)
embedding2/embeddings:0 (4, 300)
bidirectional/forward_gru/gru_cell_1/kernel:0 (300, 24)
bidirectional/forward_gru/gru_cell_1/recurrent_kernel:0 (8, 24)
bidirectional/forward_gru/gru_cell_1/bias:0 (2, 24)
bidirectional/backward_gru/gru_cell_2/kernel:0 (300, 24)
bidirectional/backward_gru/gru_cell_2/recurrent_kernel:0 (8, 24)
bidirectional/backward_gru/gru_cell_2/bias:0 (2, 24)
attention/att_weight:0 (16, 1)
attention/att_bias:0 (34, 1)
dense/kernel:0 (16, 8)
dense/bias:0 (8,)
dense_1/kernel:0 (8, 4)
dense_1/bias:0 (4,)
time_distributed/kernel:0 (4, 1)
time_distributed/bias:0 (1,)

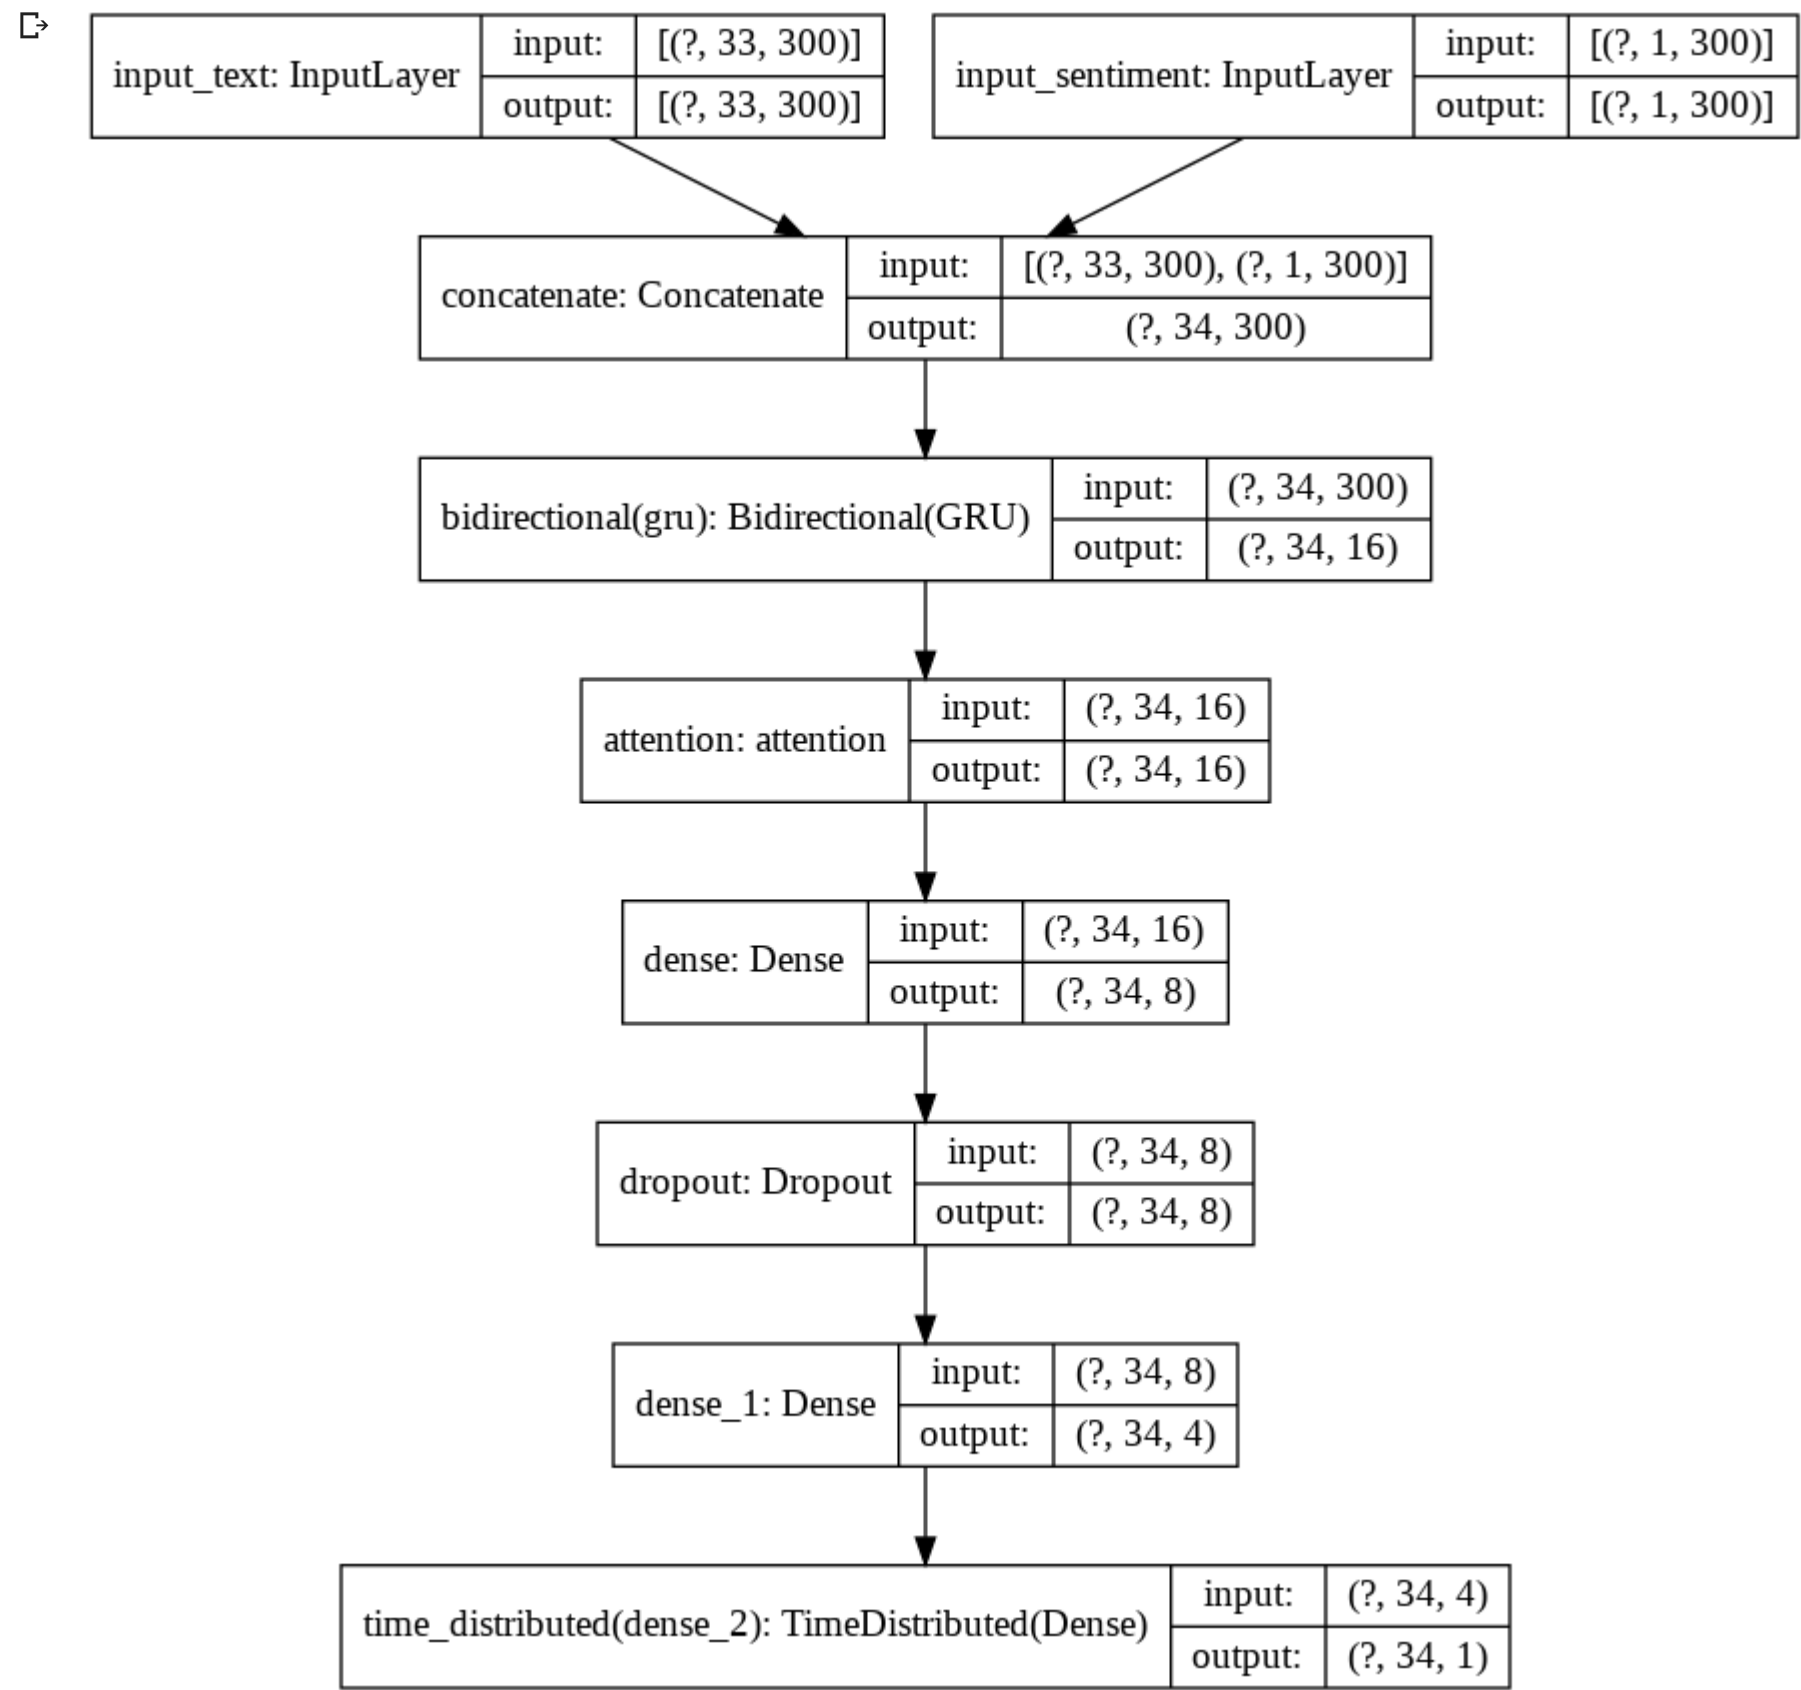
tf.keras.backend.clear_session()

from tensorflow.keras.models import Model

```
input1=Input(shape=(max_length_text,300),name='input_text')
input2=Input(shape=(max_length_sentiment,300),name='input_sentiment')

concat= Concatenate(axis=1)([input1,input2])
gru=Bidirectional(GRU(8,name='gru',return_sequences=True,dropout=0.4))(concat)
att = attention(return_sequences=True)(gru)
dense1 = Dense(8,activation='relu',kernel_regularizer=l2(0.0001))(att)
dp = Dropout(0.5)(dense1)
dense1 = Dense(4,activation='relu',kernel_regularizer=l2(0.0001))(dp)
output=TimeDistributed(Dense(1,activation='sigmoid'))(dense1)

model = Model(inputs=(input1,input2),outputs=output)
import tensorflow as tf
tf.keras.utils.plot_model(model, 'Model1.png',show_shapes=True)
```



```
model.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_text (InputLayer)	[(None, 33, 300)]	0	
input_sentiment (InputLayer)	[(None, 1, 300)]	0	
concatenate (Concatenate)	(None, 34, 300)	0	input_text[0][0] input_sentiment[0][0]
bidirectional (Bidirectional)	(None, 34, 16)	14880	concatenate[0][0]
attention (attention)	(None, 34, 16)	50	bidirectional[0][0]
dense (Dense)	(None, 34, 8)	136	attention[0][0]
dropout (Dropout)	(None, 34, 8)	0	dense[0][0]
dense_1 (Dense)	(None, 34, 4)	36	dropout[0][0]
time_distributed (TimeDistribut	(None, 34, 1)	5	dense_1[0][0]
=====			
Total params: 15,107			
Trainable params: 15,107			
Non-trainable params: 0			

```
model.set_weights(w[2:])
```

```
x_val = x_val[['text','selected_text','sentiment']]
x_val
```



	text	selected_text	sentiment
639	sorry to tweet about bgt but poor wonderful cr...	sorry	negative
9466	walah me still i am not aettina the full idea	not aettina the full idea	neative

```
from nltk.tokenize import WhitespaceTokenizer
def get_text_vectors(sent):
    vector = np.zeros(shape=(33,300))
    tk = WhitespaceTokenizer()
    sent = tk.tokenize(sent)
    for i in range(len(sent)):
        try:
            vec = fasttext_model.wv.word_vec(sent[i])
            vector[i] = vec
        except:
            continue
    return vector
```

all_sent = (x_val['text'].values)

```
ip_text = np.zeros(shape=(len(all_sent),33,300))
for i in tqdm(range(len(all_sent))):
    res = get_text_vectors(all_sent[i])
    res = res.reshape((1,res.shape[0],res.shape[1]))
    ip_text[i] = res
```

100%|██████████| 3271/3271 [00:00<00:00, 10508.97it/s]

ip_text.shape

(3271, 33, 300)

```
def get_sentiment_vector(sentiment):
    return fasttext_model.wv.word_vec(sentiment)
```

all_sentiment = x_val['sentiment'].values

```
ip_sentiment = np.zeros(shape = (len(all_sentiment),1,300))
for i in tqdm(range(len(all_sentiment))):
    vec = get_sentiment_vector(all_sentiment[i])
    vec = vec.reshape((1,1,300))
    ip_sentiment[i] = vec
```

ip_sentiment.shape

100%|██████████| 3271/3271 [00:00<00:00, 177512.27it/s]
(3271, 1, 300)

val_data = (ip_text,ip_sentiment)

```
val_pred = model.predict(val_data)
val_pred = np.squeeze(val_pred)
val_pred = np.where(val_pred>0.2,1,0)
val_pred.shape
```

(3271, 34)

```
val_pred_output = []
for each in tqdm(val_pred):
    indices=[]
    for x in range(len(each)):
        if each[x] == 1:
            indices.append(x)
        else:
            continue
```

indices = np.array(indices)
val_pred_output.append(indices)

print(len(val_pred_output))

100%|██████████| 3271/3271 [00:00<00:00, 62916.19it/s]3271

x_val['prediction'] = val_pred_output

x_val

	text	selected_text	sentiment	prediction
639	sorry to tweet about bgt but poor wonderful cr...	sorry	negative	[0, 1, 6, 7, 8, 9, 11, 12, 13, 14, 15]
9466	walah me still i am not getting the full idea	not getting the full idea	negative	[5, 6]
7105	yo wake your curse up and go to work go get th...	sick dont	negative	[0, 1, 2, 3, 13, 14, 15]
19196	thanks for sharing	thanks	positive	[0, 1, 2]
4717	it is an app to finally face the truth you lac...	you lack time	negative	[6, 7, 8, 9, 10, 11, 12, 13, 14, 17]
...
15914	oh noesss seniors last day however tickling wi...	totally worth it	positive	[0, 1, 2, 14, 15, 16, 17, 18]
23340	good luck tomorrow	good luck tomorrow	positive	[0, 1, 2]
12499	played with fontstruct links uploaded to dafon...	top	positive	[13]
21280	think i am gonna start writing a proper blog c... i am gonna start writing a proper blog		positive	[11, 12, 13, 14]
18594	not to sound preachery or anything but my ipho...	amazing	positive	[0, 1, 2, 9, 10, 11]

3271 rows × 4 columns

```
def get_pred_text(x):
    pred = []
    text = x[0].split()
    indices = x[1]
```

```
l = len(text)
for each in indices:
    if each < l:
        pred.append(text[each])

return pred
```

```
pred_text= x_val[['text', 'prediction']].progress_apply(lambda x:get_pred_text(x),axis=1)
```

```
100% 3271/3271 [00:00<00:00, 26956.14it/s]
```

```
x_val['pred_text'] = pred_text
x_val['pred_text'] = x_val['pred_text'].apply(lambda x: ' '.join(x))
x_val.sample(5)
```

	text	selected_text	sentiment	prediction	pred_text
8187	this is dumb i keep losing followers	this is dumb i keep losing followers	negative	[0, 1, 2, 3, 5]	this is dumb i losing
843	it just had to rain on me almost a perfect day...	it just had to rain on me	negative	[5, 6, 7, 8, 9, 10]	on me almost a perfect day
14300	c news was not as bad as i expected could have...	better	positive	[2, 3, 4, 5, 18]	was not as bad great
21638	omg that is awful wow our pyr figured out how ...	omg that is awful	negative	[0, 1, 2, 3, 4, 5, 17]	omg that is awful wow our crushing
8667	peep this remix from nothe wu dynasty remix ta...	delayed	negative	[10]	delayed

```
def jaccard_score(x):
    str1, str2 = str(x[0]),str(x[1])
    a = set(str1.lower().split())
    b = set(str2.lower().split())
    c = a.intersection(b)
    return float(len(c)) / (len(a) + len(b) - len(c))
```

```
x_val['jaccard'] = x_val[['selected_text', 'pred_text']].progress_apply(jaccard_score,axis=1)
```

```
100% 3271/3271 [00:00<00:00, 31845.32it/s]
```

```
x_val.sample(5)
```

	text	selected_text	sentiment	prediction	pred_text	jaccard
12935	adam samberg new moon trailor good evening too...	too bad my cable is off as of friday	negative	[5, 6, 7, 8]	good evening too bad	0.181818
21439	so jealous see if you can get some dallas conc...	jealous	negative	[0, 1, 2]	so jealous see	0.333333
11408	i might have been a child but i was never one ...	apologised	negative	[1, 2, 3, 7, 8, 9, 11, 17]	might have been i was never of apologised	0.125000
6459	do not count on it	do not count on it	negative	[1]	not	0.200000
6442	happy mother is day to all the mothers	happy mother is day	positive	[0, 1, 2, 3]	happy mother is day	1.000000

```
pos_data = x_val[x_val['sentiment'] == 'positive']
neg_data = x_val[x_val['sentiment'] == 'negative']
#neu_data = x_val[x_val['sentiment'] == 'neutral']
pos_data.shape,neg_data.shape
```

```
((1661, 6), (1610, 6))
```

Jaccard scores for validation data

```
print('Mean jaccard score for positive sentiment data:', np.mean(pos_data['jaccard']))
print('Mean jaccard score for negative sentiment data', np.mean(neg_data['jaccard']))
#print('Mean jaccard score for neutral sentiment data', np.mean(neu_data['jaccard']))
```

```
Mean jaccard score for positive sentiment data: 0.40192826596900244
Mean jaccard score for negative sentiment data 0.37148625313293876
```